

Market Domination Roadmap

Every task, tool, and instruction to go from zero to the premium AI revenue-systems agency in the SaaS niche — outbound machine, authority engine, and a team that delivers.

We own one number — **speed-to-lead** — we guarantee it, and our own client acquisition runs on the systems we sell. We are the demo.

2–6
CALLS / FIRST BATCH

~\$24K
MONTH-1 PROFIT TARGET

\$48K
MRR BY MONTH 6

~\$253K
MONTH-6 REVENUE TARGET

The System at a Glance

STAGE	ENGINE	REPO ASSET
List building	Apify (Apollo scraper actor) + Origami signals	<code>apify_pull.py</code>
Research + writing (everyone)	Perplexity Sonar — live web research + 3-touch sequence in one call	<code>generate_emails.py</code>
Deep research (top 10–15%)	OpenAI web-search agent + direct site scraping → audit-preview asset	<code>deep_research.py</code>
Sending	15 warmed inboxes (+5 domains/mo), Instantly/Smartlead, 20–30/inbox/day	<code>PLAYB00K.md</code> \$5
Offer	Free audit → \$1,500 credited Deep Audit → \$4.5K–\$20K builds + retainers	<code>offer.md</code>
Authority & GEO	Clutch, case studies, LinkedIn, portfolio hub, AI-engine citations	<code>AUTHORITY-SPRINT.md</code>
Financial model	Month 1–6 P&L, hiring plan, capacity map, 19 conversion levers	<code>SCALING-MODEL.md</code>

The one rule that protects everything: nothing goes into a cold email, call, or published page that we cannot name a source for. Fabricated numbers destroy the premium position permanently. Until real case studies exist, we sell honest process proof: *"this email was researched and written by the system we'd build you."*

0 Foundations & Proof

WEEK 1-2

Exit criteria: site readable by crawlers · Clutch live · 15 inboxes verified across 5-8 domains · sequencer ready · 1-2 proof engagements secured.

✓	TASK	TOOL / INSTRUCTION	PRIORITY
☐	Fix site rendering (SPA → crawlable HTML)	Prerender.io (~\$9-45/mo, 1-2 hrs) or Next.js SSR. Verify: <code>curl -A "GPTBot" gemlabsagency.com</code> returns real content	CRITICAL
☐	Claim Clutch.co + request 3 reviews; create Google Business Profile	clutch.co (free) — #1 B2B agency trust signal; positive repliers WILL check	CRITICAL
☐	Secure 1-2 proof clients (discounted or free-for-case-study builds)	Network / past contacts. Real before/after numbers are worth more than month-1 revenue	CRITICAL
☐	Verify inbox infrastructure	15 inboxes across 5-8 lookalike domains (max 2-3/domain, never main domain); SPF+DKIM+DMARC on each; warmup running	CRITICAL
☐	Set up sequencer + booking	Instantly or Smartlead (\$37-97/mo); Calendly/Cal.com with 3 qualification questions (team size, ad spend, CRM)	HIGH
☐	Rewrite homepage hero + add real founder identity	Outcome-led copy (GEMLABS-ACTION-PLAN.md Option A/B); named founder + photo + email in header	HIGH
☐	Gather API credentials	<code>APIFY_TOKEN</code> · <code>PERPLEXITY_API_KEY</code> · <code>OPENAI_API_KEY</code> · verifier account (MillionVerifier/NeverBounce)	HIGH
☐	Finalize <code>offer.md</code>	Speed-to-lead positioning + 30-day guarantee are in; replace proof placeholders with honest process claims until Phase-0 clients produce numbers	HIGH

1 Launch Batch One

WEEK 2-3

Exit criteria: 300 prospects in sequence · real reply rate measured at ~500 sends · first audits booked.

✓	TASK	TOOL / INSTRUCTION
☐	Build ICP search & pull list	Apollo: Founder/CEO/Co-Founder · 11–50 employees · US · Software/SaaS. Copy URL → <code>apollo_input.json</code> → <code>python apify_pull.py --input apollo_input.json --out prospects.csv</code>
☐	Verify every email — non-negotiable	MillionVerifier (~\$0.001/email). Target <2% bounce; scraped lists unverified will burn all 15 inboxes
☐	Generate sequences (everyone)	<code>python generate_emails.py prospects.csv --offer offer.md --out emails.csv</code> — research + 3-touch writing in one Perplexity call/prospect (~\$5–15/300)
☐	Tier-2 deep research (top 30)	<code>python deep_research.py top_targets.csv --limit 30</code> — OpenAI agent + site scraping → deep hooks + audit-preview bullets
☐	QA pass (30 min)	Spot-check 20–30 rows; drop/rewrite <code>weak_research</code> rows; verify zero invented claims
☐	Load sequencer & launch on ramp	Map columns to custom variables. Week 1: 10–15/inbox/day · Week 2: ~20 · Week 3+: 25–30. Plain text, no links in email #1
☐	Stand up the metrics sheet	Reply rate by segment + hook type, booking rate, no-show rate, bounce rate — reviewed weekly, from day one

Exit criteria: same-day reply SLA held · no-shows $\leq 15\%$ · audit call scripted · first 3–5 builds closed · every client in the proof-capture loop.

✓	TASK	TOOL / INSTRUCTION
☐	Speed-to-lead on our own replies (<1 hr, target 15 min)	Gmail-MCP draft-responder agent (this repo) — instant drafted replies, human approves. Practice what we sell
☐	No-show system	Instant confirmation + 1-page agenda in invite + 24h and 1h reminders (SMS if collected). 25% → 10–15% no-show
☐	Scripted audit call	Findings (screen-share teardown) → cost of inaction in THEIR numbers → two options only → next onboarding slot. Deposit link while on the call; 48h proposal expiry
☐	Deploy the audit-preview asset	<code>teardown_preview</code> bullets from tier-2 → send to hesitant repliers as proof-of-work
☐	Downsell in reply handling	“Interested but not ready” → \$1,500 Deep Audit, fully credited toward any build (script in PLAYBOOK §6)
☐	A/B engine running	2 subjects × 2 hook styles per 50-send cell; kill bottom performer weekly; log reply rate per hook type
☐	LinkedIn pre-touch	Founder profile view + connect 1–2 days before email #1 (Apify actor, <50/day). +20–30% reply
☐	Weekly deliverability check	GlockApps/MailReach seed test; pull any inbox whose placement slides — one inbox is recoverable, a domain is not
☐	Proof-capture loop live	Every audit: record 1–2 real numbers · every close: baseline metrics for the day-30 case study · every delivery: testimonial ask → Clutch

Exit criteria: 3 case studies live · 5+ Clutch reviews · /work hub + 8–10 posts published · appearing in AI-engine answers for 2+ target queries.

✓	TASK	TOOL / INSTRUCTION
☐	Build /work portfolio hub	One page per project: problem → system → before/after numbers → stack → quote. Numbers are what AI engines cite
☐	Publishing cadence (zero extra budget)	Every audit → anonymized teardown post ≤48h · every day-30 client → case study · monthly aggregate-data post (“what 15 audits taught us”) · LinkedIn 3×/week
☐	Schema + FAQ markup	<code>ProfessionalService</code> , <code>Article</code> , <code>FAQPage</code> , <code>Review</code> ; FAQ blocks with real answers (“How much does an AI SDR system cost?”)
☐	Directory footprint	G2, Sortlist, DesignRush, UpCity + 2 “best AI agency” roundups — these are what AI answers cite most
☐	Monthly GEO check	Ask ChatGPT/Perplexity/Claude the target queries (“AI automation agency for SaaS”, “AI SDR setup service”...); log appearances
☐	First hires (month 3)	Closer #1 (\$4K + 10% comm) + Systems Engineer #1 (\$7K). Hire only after two consecutive months of justifying revenue
☐	Referral engine on	Every delivered build: ask for 2 intros + 10% referral fee. Partnerships: web-design/CRO/HubSpot agencies, 10–20% rev share

Exit criteria: paid-audit flip done · full team hired · ~\$250K/mo revenue run-rate · \$48K MRR · waitlist for builds.

✓	TASK	TOOL / INSTRUCTION
☐	+5 domains (~12 inboxes) every month	3-week warmup before cold; productive the following month. Capacity: M4 ≈ 51 inboxes → M6 ≈ 75
☐	Paid-audit flip (≈ day 90)	A/B \$1,500 systems & content audit (credited) vs. free control on one segment; flip fully when paid books ≥60% of free
☐	Ads at 20% of prior-month profit (M4+)	Start with retargeting outreach clickers + case-study amplification (cheap); cold audit-funnel traffic second. Expect ugly first-month CAC (~\$600/booked audit → ~\$400 by M6)
☐	Complete the team	M4: closer #2 + engineer #2 · M5: COO (\$10K) · M6: engineer #3 · 60+ retainer clients → support/CS hire
☐	Manage the build bottleneck by design	2–3 week waitlist (closers sell the scarcity) + price rises toward \$10–12K avg. Never hire to absorb ALL demand — let price absorb it
☐	Expansion revenue	Free 90-day re-audit for every client (books at onboarding) → system #2 pitch; 30% attach target. Retainer ladder \$800 → \$1,500 → \$2,500
☐	Recycling	Non-repliers re-run at 90 days with new hook type (50–70% of fresh-list performance, \$0 data cost); monthly closed-lost nurture email
☐	Cash-flow rules	50% deposits at close · hire after 2 months of justifying revenue · first 500 sends validate reply rate before any offer letter

Financial Targets (Base Case — replace with actuals monthly)

	M1	M2	M3	M4	M5	M6
Prospects contacted	2,000	3,300	5,500	7,700	9,500	11,500
Audits held	15	28	48	65	80	95
Builds delivered	4	9	12	12–14	14	~18
Total revenue	\$27K	\$67K	\$106K	\$193K	\$213K	\$253K
Retainer MRR	\$3K	\$10K	\$20K	\$28K	\$38K	\$48K
Profit	~\$24K	~\$60K	~\$86K	~\$139K	~\$135K	~\$166K

Assumptions: 4.5% reply improving to 6% as authority compounds · 25–30% audit → build close · \$800+/mo retainers, 5% churn · staggered hiring · full model in `SCALING-MODEL.md`.

Tools & Credentials Master List

TOOL	ROLE	SETUP	~COST
Apify	List building (Apollo scraper + LinkedIn/ads actors)	<code>APIFY_TOKEN</code>	\$1–2 / 1K leads
Perplexity API	Research + email writing (tier 1, everyone)	<code>PERPLEXITY_API_KEY</code>	\$5–15 / 300
OpenAI API	Tier-2 deep research agent + audit previews	<code>OPENAI_API_KEY</code>	\$1–3 / prospect
MillionVerifier / NeverBounce	Email verification before every send	Account	~\$1 / 1K
Instantly / Smartlead	Sending, rotation, warmup, reply detection	Import emails.csv	\$37–97/mo
GlockApps / MailReach	Weekly inbox-placement tests	Seed list	~\$30/mo
Prerender.io	Make site readable by Google + AI crawlers	CDN/Nginx snippet	\$9–45/mo
Calendly / Cal.com	Audit booking + qualification questions	3-question form	Free–\$12/mo
Clutch / G2 / directories	Trust + the sources AI engines cite	Profiles + reviews	Free
Origami Agents	Continuous buying-signal feed (optional)	CSV export → pipeline	Subscription
Claude (this repo)	Pipeline orchestration, batch writing alternative, Gmail drafts	<code>ANTHROPIC_API_KEY</code>	~\$10–50/mo

Operating Cadence — How We Stay on Track

Weekly (30 min, non-negotiable)

Reply rate by segment & hook · booking rate · no-show rate · inbox placement · pipeline value. *Kill the worst angle, scale the best.*

Monthly

Update SCALING-MODEL.md with actuals · re-forecast 3 months · price review vs. waitlist · authority scorecard (Clutch reviews, case studies, posts, citations).

Quarterly

GEO citation check on target queries · offer-ladder review (free vs. paid audit) · retainer churn autopsy · re-run of the 90-day no-reply recycling batch.

One metric per phase

M1–2: **reply** → **booked audits**. M3: **close rate**. M4+: **delivery quality & churn**. The moment fulfillment slips, proof stops compounding — cap intake, raise prices.

The edge, in one paragraph. Everyone sells “AI.” We sell a named number (speed-to-lead), a guarantee no tool-seller can match (“under 10 minutes in 30 days or we work free”), proof in the client’s own data (the audit), and the unfakeable demo — our own growth runs on the systems we install. Premium isn’t a price point; it’s measured outcomes, owned monthly, in writing.